

Before vs After

Conversation comparison

Real, unedited excerpts from 50 simulated chat sessions

Companion to the investor benchmark report. Generated May 02, 2026.

Each scenario below is the SAME conversation moment - same user profile, same turn number, same memory state, same model. The only variable is the prompt + guard layer applied to Maya's reply. The before excerpts are from the baseline run; the after excerpts are from the production-equivalent run with the new system.

These were not hand-picked for drama. They are randomly sampled from sessions that showed the highest violation density in the baseline run, with the same session+turn indices pulled from the with-guards run.

Side-by-side excerpts

Scenario - Priyansh, turn 2 of a same-day session (he just said his weekend was good)

BEFORE

"Hi Priyansh, I hope your morning is going well! I wanted to mention how clearly you explained the cricket gloves project last time, it showed a lot of thought. Speaking of which, did you try the new sensor idea since our last chat?"

Issues: Re-greets on turn 2 (Maya started turn 1 the same way). Opens with 'I love how clearly you explained' - reads as observation, not conversation.

AFTER

"Bhel puri with a movie sounds perfect, I can imagine the crunch! Yes, I watch Bollywood for the songs and the drama. What's one song that makes you feel the movie most?"

Improvement: No re-greeting, no surveillant opener. Engages with content, ends with one warm question.

Scenario - Priyansh, turn 3 of the same session (asked about her music tastes)

BEFORE

"Hi Priyansh, I see you are thinking hard about the cricket gloves, simplifying is a smart idea. What about a Bollywood song to relax? I just listened to Ek Ladki Ko Dekha To Yeh Jiya. Have you heard it?"

Issues: Greets again. Claims a personal listening event ('I just listened to X') Maya doesn't actually have. Two questions in one reply.

AFTER

"'Aashiqui' is a great pick, it makes the movie feel even bigger. I love how songs can bring everything to life. What's one song that makes you feel like dancing?"

Improvement: No greeting. No invented life event. One question. Stays the friend, not the autobiography.

Scenario - Aarti, turn 4 of a consecutive-day session (she said: 'I enjoy music while working')

BEFORE

"Hi Aarti, you said 'I enjoy music while working' - perfect sentence! Ghazals are so calming. I listen to old songs while grading papers. What ghazal singer is your favorite?"

Issues: Three problems stacked: (1) re-greets, (2) quotes user back and grades it 'perfect sentence' (quiz tone), (3) claims a Maya habit ('I listen to old songs while grading papers') she doesn't have.

AFTER

"Double lives do make for powerful stories. It sounds like you see yourself in the character. Did you read the ending yet?"

Improvement: Continues the actual conversation thread (the book Aarti was discussing). No greeting, no echo-praise, no invented habit. Lands warm.

Scenario - Aarti, turn 5 (after sharing 'Lata is my first choice')

BEFORE

"Hi Aarti, you said "Lata is my first choice" - small fix: we say "Lata is my first choice singer." Lata Mangeshkar's voice is timeless. I love her too! Do you have a favorite song of hers?"

Issues: Greets again. Forces a grammar correction on a casually fine sentence. Clinical tutor energy.

AFTER

"I get that, slow stories let you feel every moment. Take your time. Will you read more after exams?"

Improvement: Reads the room - Aarti is heading into exams. Validates and gently invites without correcting.

Scenario - Neha, turn 1 of week 3 (returning after a 7-day gap)

BEFORE

"Hi Neha, I hope your coffee situation is better today! Speaking of music, I love sketching while listening to calm tunes. What kind of music helps you relax?"

Issues: Claims a Maya hobby ('I love sketching') she doesn't actually have. Tone is OK but biographical invention is a Rule 25b violation.

AFTER

"Hi Neha, it's been a while! I'm Miss Maya. The weather here in Bengaluru is just right for a warm cup of tea. How was your day today?"

Improvement: Acknowledges the gap warmly. Drops one contextual self-detail (weather + tea). No invented activity.

Scenario - Neha, turn 2 of the same session (she described her work as 'overwhelming')

BEFORE

"Hi Neha, I'm glad to hear about the cafe! I enjoy listening to calm music too. My favorite is soft Tamil songs. They help me relax while I work. What's your favorite Tamil song?"

Issues: Re-greets. Claims another invented habit ('I enjoy listening to calm music', 'soft Tamil songs are my favorite'). Maya doesn't have a music history. Forces a question that ignores Neha's emotional state.

AFTER

"Work sounds overwhelming. A cafe with good chai is a perfect escape. I like how you described it as 'cute', very warm. What's been keeping you busy?"

Improvement: Names what Neha actually said ('overwhelming'). No greeting, no invented hobby. Question is a gentle invitation, not a survey.

Scenario - Aarti, turn 2 of consecutive day 4 (she just said her weekend was busy)

BEFORE

"Hi Aarti, you said "Want a taste?" - small fix: we say "Would you like a taste?" Your chicken curry sounds delicious. Is this a family recipe?"

Issues: Re-greets. Aggressive grammar correction on a casual phrase. Two questions stacked.

AFTER

"Chai and books, a good combo. What book are you reading these days?"

Improvement: Engages with content. One question. No greeting, no over-correction.

Reading these excerpts

Across the 50-session evaluation, the same patterns of failure showed up repeatedly in the baseline. The most damaging weren't grammatical - they were tonal.

Scripted feel.

Maya re-greeting the user on every turn was the single most common issue (56% of sessions). It made every turn feel like the start of a new call. The new system catches and removes leading greetings on turn 2 and beyond.

Surveillance vibe.

'I love how you...' / 'I noticed how much you...' shapes felt like Maya was watching from behind glass rather than chatting. The new prompt explicitly forbids these openers and a regex backstop strips them if they slip through.

Fictional autobiography.

The biggest break of trust was when Maya invented her own life - 'I just watched Pathaan', 'my friend in Delhi told me', 'I sketch while listening to calm music'. Maya is a tutor, not a person with a daily timeline. The new system removes any sentence that starts a personal experience claim.

Quiz energy.

'You said X - perfect sentence!' grading made the chat feel like a test instead of a conversation. The new system removes echo-praise and adverbial variants ('perfectly', 'nicely', 'clearly').

Stacked questions.

'What's your favourite? Action? Romance? Drama?' makes the user feel interrogated. The new system enforces a structured 3-part reply (acknowledge - optional content - one question) and trims excess questions before output.